



# Concurrency Control

- ➔ • Schedule
  - Conflict Serializable Schedule
  - Two phase locking
  - Warning Protocol
  - Validation

Database System - Nankai



## Week15\_Course

### Database System\_Concurrency Control part1


#### Schedule

Database System - Nankai



# Concurrency Control

T1 T2 ... Tn



How to prevent harmful interference between transactions?

=> **scheduling techniques** based on

- locks
- timestamps and validation

Database System - Nankai



## Correctness depends on scheduling of transactions

### A schedule

- Chronological (possibly interleaving) order in which actions of transactions are executed
- A correct schedule is equivalent to executing transactions one-at-a-time in some order

T1	T2	T3
Write (A)	Write (B)	Write (C)
Write (B)	Write (C)	Write (A)
Write (D)	Write (D)	

Database System - Nankai



## Example:

Constraint:  $A=B$

<p>T1: Read(A)  <math>A \leftarrow A+100</math>  Write(A)  Read(B)  <math>B \leftarrow B+100</math>  Write(B)</p>	<p>T2: Read(A)  <math>A \leftarrow A \times 2</math>  Write(A)  Read(B)  <math>B \leftarrow B \times 2</math>  Write(B)</p>
---	---

Database System - Nankai



## Schedule A

T1	T2	A	B
Read(A); $A \leftarrow A+100$ ;		25	25
Write(A);			
Read(B); $B \leftarrow B+100$ ;		125	
Write(B);			125
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	250	
	Read(B); $B \leftarrow B \times 2$ ;		
	Write(B);		250
<div>OK</div> <div> <math>S_A = r1(A)w1(A) r1(B)w1(B)r2(A)w2(A)r2(B)w2(B)</math> </div>		250	250

Database System - Nankai



## Schedule B

T1	T2	A	B
		25	25
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	50	
	Read(B); $B \leftarrow B \times 2$ ;		50
	Write(B);		
Read(A); $A \leftarrow A + 100$ ;		150	
Write(A);			
Read(B); $B \leftarrow B + 100$ ;			150
Write(B);			
		150	150

OK

Database System - Nankai




## Schedule C

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);			
	Read(A); $A \leftarrow A \times 2$ ;	125	
	Write(A);		
Read(B); $B \leftarrow B + 100$ ;			125
Write(B);			
	Read(B); $B \leftarrow B \times 2$ ;		250
	Write(B);		
		250	250

OK

$Sc = r1(A)w1(A)r2(A)w2(A)r1(B)w1(B)r2(B)w2(B)$




Database System - Nankai



## Schedule D

T1  
Read(A);  $A \leftarrow A+100$ ;  
Write(A);



Read(B);  $B \leftarrow B+100$ ;  
Write(B);

T2  
Read(A);  $A \leftarrow A \times 2$ ;  
Write(A);  
Read(B);  $B \leftarrow B \times 2$ ;  
Write(B);

Constraint violation!

A	B
25	25
125	
250	
	50
	150
250	150

Database System - Nankai



## Schedule E

Same as Schedule D  
but with new T2'

T1  
Read(A);  $A \leftarrow A+100$ ;  
Write(A);

Read(B);  $B \leftarrow B+100$ ;  
Write(B);

T2'  
Read(A);  $A \leftarrow A \times 1$ ;  
Write(A);  
Read(B);  $B \leftarrow B \times 1$ ;  
Write(B);

A	B
25	25
125	
125	
	25
	125
125	125

Database System - Nankai



- Want schedules that are “good”, regardless of
  - \_ initial state ( $\leftrightarrow$  “good” in any DB state) and
  - \_ transaction semantics
- Only look at order of READs and WRITEs
  - \_ Note: transactions see values in buffers, not on disk
  - => this time ignore INPUT/OUTPUTs

Example:  $S_a$  (Schedule a) =

$\underbrace{r1(A)w1(A)r1(B)w1(B)}_{T1} \quad \underbrace{r2(A)w2(A)r2(B)w2(B)}_{T2}$

Database System - Nankai



## Scheduling Transactions: Definitions

- Serial schedule: no concurrency
  - \_ Does not interleave the actions of different transactions.
- Equivalent schedules: same result on any DB state
  - \_ For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- Serializable schedule: equivalent to a serial schedule
  - \_ A schedule that is equivalent to *some* serial execution of the transactions.
  - (Note: If each transaction preserves consistency, every serializable schedule preserves consistency. )

Database System - Nankai



# Serial Schedule


- A schedule is **serial**, if actions of transactions are not interleaved
  - \_ e.g., (T1, T2) or (T2, T1)
  - \_ A serial schedule obviously maintains consistency (assuming correctness of individual transactions)
- Could we reorder a schedule into an equivalent serial schedule?
  - \_ Actions **conflict**, if swapping them may change the meaning of a schedule:
    - any two actions of a single transaction
    - two actions on a common DB element A, one of which is WRITE(A)

Database System - Nankai



Example: (of swapping non-conflicting actions)

$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$





Database System - Nankai



However, for  $S_d$ :

$S_d = r_1(A)w_1(A)r_2(A)w_2(A)r_2(B)w_2(B)r_1(B)w_1(B)$



- $S_d$  cannot be rearranged into a serial schedule

⇒  $S_d$  is not "equivalent" to any serial schedule

⇒  $S_d$  is "bad"

Database System - Nankai



## Summary

- Schedule
- Serial Schedule

Database System - Nankai



多选题 1分



### 互动交流一 —— 不定项选择题

以下哪一种写法表示的是事务的调度 (Schedule) ?

- ☐ A  $T1 = R1(C) R1(A) W1(C) R1(B) W1(D) W1(B) R1(A)$
- ☐ B  $S1 = W3(A) R1(B) R2(A) R4(D) W1(B) R4(A) W2(B) W1(D) R4(B)$
- ☐ C  $T2 = R1(C) W1(D) W1(B) \quad T3 = R2(A) W2(D) W2(B)$
- ☐ D  $S2 = W3(A) R1(B) W1(D) W1(B) R2(A) W2(B) R4(D) R4(A) R4(B)$

提交

Database System - Nankai

多选题 1分



### 互动交流二

以下哪一个调度是串行调度 (Serial Schedule) ?

- ☐ A  $T1 = R1(C) R1(A) W1(C) R1(B) W1(D) W1(B) R1(A)$
- ☐ B  $S1 = W3(A) R1(B) R2(A) R4(D) W1(B) R4(A) W2(B) W1(D) R4(B)$
- ☐ C  $T2 = R1(C) W1(D) W1(B) \quad T23 = R2(A) W2(D) W2(B)$
- ☐ D  $S2 = W3(A) R1(B) W1(D) W1(B) R2(A) W2(B) R4(D) R4(A) R4(B)$

提交

Database System - Nankai



## Week15\_Course

### Database System\_Concurrency Control part2

#### Conflict Serializable Schedule

Database System - Nankai



## Concurrency Control

- Schedule
- ➔ • Conflict Serializable Schedule
- Two phase locking
- Warning Protocol
- Validation

Database System - Nankai



## Concepts

**Transaction:** sequence of  $r_i(x)$ ,  $w_i(x)$  actions

**Conflicting actions:**

$$\begin{array}{ccc} r_h(A) & w_h(A) & w_h(A) \\ < w_k(A) < r_k(A) < w_k(A) \end{array}$$

If schedule  $S$  contains conflicting actions

$\dots, p_h(A), \dots, q_k(A), \dots$  [i.e., one of  $p, q$  is  $w$ ],

transaction  $T_h$  must precede  $T_k$  in a corresponding serial schedule. Denote this by  $T_h \rightarrow T_k$

Database System - Nankai



## Returning to Sc

$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$

$\swarrow \quad \searrow \quad \quad \quad \swarrow \quad \searrow$   
 $T_1 \rightarrow T_2 \quad \quad \quad T_1 \rightarrow T_2$

No cycles  $\Rightarrow Sc$  is "equivalent" to a serial schedule  
(in this case  $T_1, T_2$ )

Database System - Nankai




## Definition

$S_1, S_2$  are **conflict equivalent** schedules  
if  $S_1$  can be transformed into  $S_2$  by a series  
of swaps on non-conflicting actions.

( $\Rightarrow$  effect of both  $S_1$  and  $S_2$  on the DB is the same)

A schedule is **conflict serializable** if it  
is conflict equivalent to some serial  
schedule.



Database System - Nankai



## Definition

A schedule is **conflict serializable** if it is  
conflict equivalent to some serial  
schedule.

**NB:** Conflict serializability is a sufficient (but not  
a necessary) condition for serializability  
(equivalence to some serial schedule)  
Easier to enforce than serializability, therefore  
generally assured by commercial systems

Database System - Nankai



## Precedence graph $P(S)$ ( $S$ is schedule)

Nodes: transactions  $T_1, T_2, \dots$  in  $S$

Arcs:  $T_i \rightarrow T_j$  for  $i \neq j$  whenever

- $p_i(A), q_j(A)$  are conflicting actions in  $S$ ,  
(same element  $A$ , at least one of actions is a write)
- action  $p_i(A)$  precedes  $q_j(A)$  in  $S$

- What is  $P(S)$  for  
 $S = w_3(A) w_2(C) r_1(A) w_1(B) r_1(C) w_2(A) r_4(A) w_4(D)$
- Is  $S$  serializable?

Database System - Nankai



## Exercise:

- What is  $P(S)$  for  
 $S = w_3(A) w_2(C) r_1(A) w_1(B) r_1(C) w_2(A) r_4(A) w_4(D)$
- Is  $S$  serializable?

Database System - Nankai



Lemma Let  $S_1, S_2$  be schedules for the same set of transactions

$S_1, S_2$  conflict equivalent  $\Rightarrow P(S_1) = P(S_2)$

Proof:

Assume  $P(S_1) \neq P(S_2)$

$\Rightarrow \exists T_i: T_i \rightarrow T_j$  in  $P(S_1)$  and not in  $P(S_2)$

$\Rightarrow S_1 = \dots p_i(A) \dots q_j(A) \dots$	$\left\{ \begin{array}{l} p_i, q_j \\ \text{conflict} \end{array} \right.$
$S_2 = \dots q_j(A) \dots p_i(A) \dots$	

$\Rightarrow S_1, S_2$  not conflict equivalent

Database System - Nankai



Note:  $P(S_1) = P(S_2) \not\Rightarrow S_1, S_2$  conflict equivalent

Counter example:

$S_1 = w_1(A) \ r_2(A) \ w_2(B) \ r_1(B)$

$S_2 = r_2(A) \ w_1(A) \ r_1(B) \ w_2(B)$

Database System - Nankai



## Theorem I

$P(S_1)$  acyclic  $\iff S_1$  conflict serializable

( $\Leftarrow$ ) Assume  $S_1$  is conflict serializable

$\Rightarrow \exists$  serial  $S_s$ :  $S_s, S_1$  conflict equivalent

$\Rightarrow P(S_s) = P(S_1)$  [ $\Leftarrow$  Lemma]

$\Rightarrow P(S_1)$  acyclic since  $P(S_s)$  is acyclic

Database System - Nankai



## Theorem (cont.)

$P(S_1)$  acyclic  $\iff S_1$  conflict serializable

( $\Rightarrow$ ) Assume  $P(S_1)$  is acyclic

Transform  $S_1$  as follows:

(1) Take  $T_1$  to be transaction with no incoming arcs


(2) Move all  $T_1$  actions to the front

$S_1 = \dots q_j(X) \dots p_1(A) \dots$



(3) we now have  $S_1 = \langle T_1 \text{ actions} \rangle \langle \dots \text{rest} \dots \rangle$

(4) repeat above steps to serialize rest!



Database System - Nankai



## Example of Theorem Application

For each of the following schedules, answer the questions below:

$S_b = W3(\text{A})R1(\text{B})R2(\text{A})R1(\text{D})W1(\text{B})R4(\text{A})W2(\text{B})R4(\text{D})R3(\text{D})R4(\text{B})$

- (a) What is the precedence graph for the schedule  $S_b$ ?
- (b) Is the schedule conflict serializable? If so, show all equivalent serial transaction orders. If not, describe why not.

Database System - Nankai



## Summary

- Schedule, Serial Schedule
- Conflict Serializable Schedule
- Precedence Graph  $P(S)$
- $P(S_1)$  acyclic  $\iff S_1$  conflict serializable

Database System - Nankai



多选题 1分



### 互动交流一——不定项选择题

以下哪一个调度可以保证数据库的一致性？

A

T1:  $A=A+100$ ,  $B=B-100$   
T2:  $A=1.06*A$ ,  $B=1.06*B$

B

T1:  $A=A+100$ ,  $B=B-100$   
T2:  $A=1.06*A$ ,  $B=1.06*B$

C

T1:  $B=B-100$ ,  $A=A+100$   
T2:  $B=1.06*B$ ,  $A=1.06*A$

提交

Database System - Nankai

多选题 1分



### 互动交流二——不定项选择题

以下调度的哪个部分应该在优先图中出现？

$S = W3(\text{A}) R1(\text{B}) R2(\text{A}) R1(\text{D}) R2(\text{B}) W4(\text{A}) R1(\text{A})$

A

$T3 \rightarrow T2 \rightarrow T4 \rightarrow T1$

B

$T1 \rightarrow T2 \rightarrow T4$

C

$T1 \rightarrow T3$

提交

Database System - Nankai



## Week15\_Course

### Database System Concurrency Control part3

#### Two phase locking

Database System - Nankai



## Concurrency Control

- Schedule
- Conflict Serializable Schedule
- ➔ • Two phase locking
- Warning Protocol
- Validation

Database System - Nankai



## How to enforce serializable schedules?

### Option 1: (Optimistic strategy)

Run system, recording P(S); At end of day, check P(S) for cycles, and declare if execution was good

### Option 2: (Pessimistic strategy)

Prevent occurrence of cycles in P(S)


Database System - Nankai



## How to enforce serializable schedules?

### Option 2: (Pessimistic strategy)

Prevent occurrence of cycles in P(S)



Database System - Nankai




## A locking protocol

Two new actions:

lock (exclusive):  $l_j(A)$

unlock:  $u_j(A)$



Database System - Nankai



## Rule #1: Well-formed transactions

$T_i: \dots l_i(A) \dots p_i(A) \dots u_i(A) \dots$

- Lock elements (A) before accessing them  
( $p_i$  is a read or a write)
- Eventually, release the locks ( $u_i(A)$ )

Database System - Nankai



## Rule #2 Legal scheduler

$S = \dots \dots l_i(A) \dots \dots u_i(A) \dots \dots$

$\longleftrightarrow$   
 no  $l_j(A)$  for  $i \neq j$

- At most one transaction  $T_i$  can hold a lock on any element  $A$

Database System - Nankai



## Exercise:

- What schedules are legal?  
What transactions are well-formed?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$   
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

Database System - Nankai



## Schedule F (with simple locking)

		A	B
T1	T2	25	25
$l_1(A); \text{Read}(A)$		125	
$A := A + 100; \text{Write}(A); u_1(A)$	$l_2(A); \text{Read}(A)$	250	
	$A := A \times 2; \text{Write}(A); u_2(A)$		
	$l_2(B); \text{Read}(B)$		50
$l_1(B); \text{Read}(B)$	$B := B \times 2; \text{Write}(B); u_2(B)$		150
$B := B + 100; \text{Write}(B); u_1(B)$		250	150

Constraint violation!

Database System - Nankai



Simple-minded locking not sufficient to ensure serializability (i.e., correctness)!

→ More advanced protocol known as "two phase locking"

Database System - Nankai



## Rule #3 Two phase locking (2PL)

for transactions

$T_i = \dots \dots \dots l_i(A) \dots \dots \dots u_i(A) \dots \dots \dots$




- All lock requests of a transaction have to precede its unlock requests

Database System - Nankai



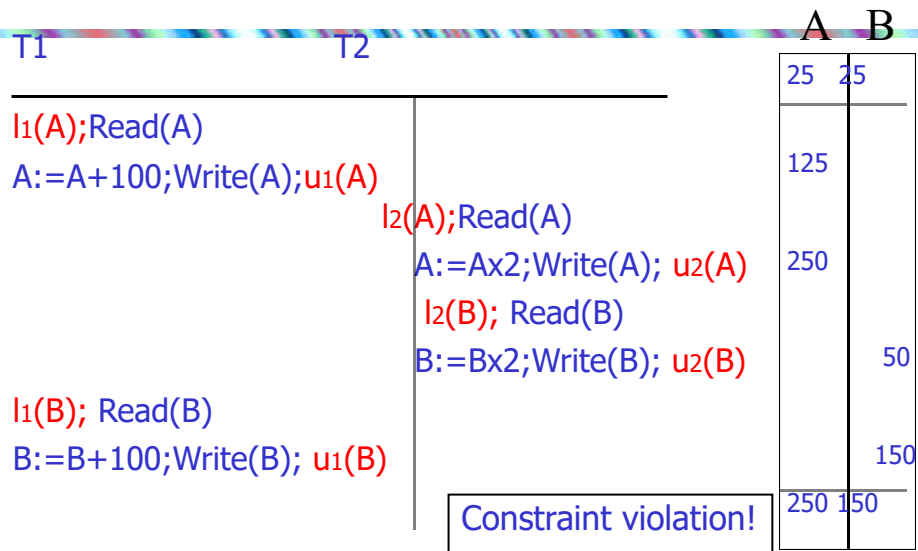
# locks  
held by  
 $T_i$



Database System - Nankai



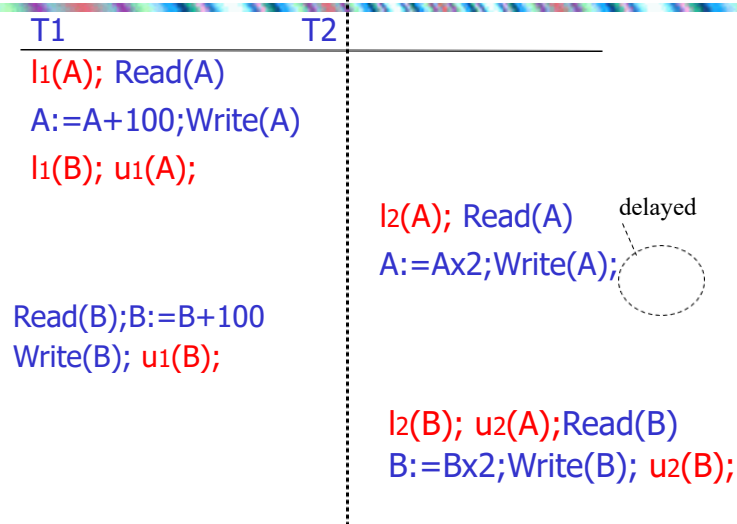
## Schedule F (with simple locking)



Database System - Nankai



## Schedule G (with 2PL)



Database System - Nankai





## Next step:


Show that Rules #1,2,3  $\Rightarrow$  conflict  
(2PL) serializable  
schedule

Database System - Nankai



To help in proof:

Definition  $\text{Shrink}(T_i) = \text{SH}(T_i) = \text{first unlock action of } T_i$



Database System - Nankai



**Lemma** Let  $S$  be a 2PL schedule.  
 $T_i \rightarrow T_j$  in  $P(S) \Rightarrow SH(T_i) <_S SH(T_j)$

Proof of lemma:

$T_i \rightarrow T_j$  means that

$S = \dots p_i(A) \dots q_j(A) \dots$ ;  $p, q$  conflict

By rules 1,2:

$S = \dots p_i(A) \dots u_i(A) \dots l_j(A) \dots q_j(A) \dots$



By rule 3:  $SH(T_i)$   $SH(T_j)$

So,  $SH(T_i) <_S SH(T_j)$

Database System - Nankai



**Theorem II Rules #1,2,3**  
**(that is, 2PL)  $\Rightarrow$  { conflict serializable schedule**

Proof: Let  $S$  be a 2PL schedule.

Assume  $P(S)$  has cycle

$T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_n \rightarrow T_1$

By Lemma:  $SH(T_1) < SH(T_2) < \dots < SH(T_1)$


Impossible, so  $P(S)$  acyclic

$\Rightarrow S$  is conflict serializable (by Th. I)

Database System - Nankai



## Schedule H (T2 reversed)



- Neither proceeds: a **deadlock**
  - System must rollback (= abort & restart) at least one of T1, T2

Database System - Nankai



## 死锁的预防

- 一次封锁法
  - 要求每个事务必须一次将所有要使用的数据全部加锁，否则就不能继续执行
  - 存在的问题：降低系统并发度、难于事先精确确定封锁对象
- 顺序封锁法
  - 顺序封锁法是预先对数据对象规定一个封锁顺序，所有事务都按这个顺序实行封锁。
  - 顺序封锁法存在的问题：维护成本高、难以实现

Database System - Nankai



## Summary-Two phase locking

- 能够判断一个事务是否是Well-formed transactions
- 能够判断一个调度是否是Legal scheduler
- 能够判断一个事务是否满足two phase locking
- 可以在一个事务中添加加锁和解锁动作，使其满足规则1和规则3
- 对满足规则1和规则3的事务可以生成可串行化的调度

Database System - Nankai